

Job Market Analysis & Recommendation System

In this project, I developed a Job Market Analysis and Recommender System to identify trending job roles and recommend suitable positions based on user input. I began by collecting and preprocessing job data, standardizing fields like location, experience, and job type. I performed exploratory analysis to uncover patterns in demand across companies, countries, and time.

Using engineered temporal and categorical features, I applied time series analysis to detect emerging job trends and used KMeans clustering to group similar roles based on keywords and skill patterns. For personalized recommendations, I combined TF-IDF vectorization with a Nearest Neighbors model to match user descriptions to the most relevant job postings.

I built an interactive Streamlit app that accepts job descriptions, allows filtering by country, experience, and job type, and returns the top 10 recommended roles with similarity scores and keyword highlights. The app also includes visual insights and a feature to download results as a PDF.

This end-to-end project brought together data preprocessing, trend analysis, clustering, NLP, and deployment into a real-time, user-friendly job role recommender.



by **Debasis Baidya**



Task 1: EDA on Upwork Job Data

Exploratory Data Analysis

Analyze the most in-demand skills across different job categories

```
1 from collections import Counter
2 # Extract keywords from job title
3 all_skills = []
4 for title in zip(df['title']):
5     skills = title
6     all_skills.extend(skills)
7
8 # Count occurrences of each skill
9 skill_counts = Counter(all_skills)
10
11 print('Most in-demand skills:')
12 print(f'{"Skill":>35} | {"Count":>5}')
13 print('-' * 45)
14 for skill, count in skill_counts.most_common(10):
15     print(f'{skill:35} | {count:5} times')
```

Skill	Count
Social Media Manager	419 times
Virtual Assistant	339 times
Logo Design	311 times
Video Editor	298 times
Graphic Designer	292 times
Logo design	173 times
Logo Designer	142 times
Full Stack Developer	136 times
Website Development	129 times
Appointment Setter	125 times


Identify countries with highest number of job posting

```
1 # Show all rows
2 pd.set_option('display.max_rows', None)
3
4 # Finding countries with highest number of dataset
5 country_counts = df['country'].value_counts()
6
7 print("Top 20 countries with highest number of job postings:\n")
8 print(country_counts.head(20))
9
10 # Reset to default after printing
11 pd.reset_option('display.max_rows')
```

Top 20 countries with highest number of job postings:

country	
United States	99834
United Kingdom	19149
India	15835
Australia	12618
Canada	11658
Pakistan	5291
Germany	4844
Netherlands	4435
United Arab Emirates	4039
France	3163
Spain	2519
Singapore	2501
Philippines	2426
China	2157
Israel	2082
Ukraine	2062
Italy	1791
Switzerland	1767
Bangladesh	1767
Saudi Arabia	1706

Name: count, dtype: int64



Dataset Overview

- Worked with 244,828 job postings from Upwork (Feb–Mar 2024)
- Cleaned missing data in fields like **budget** and **hourly rate**
- Converted **published_date** to datetime format for trend analysis



Skill Keyword Extraction

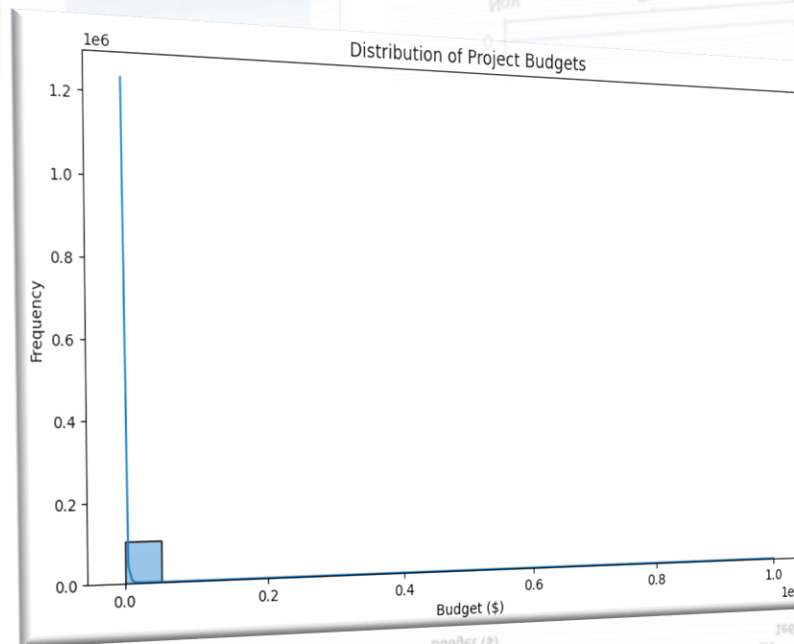
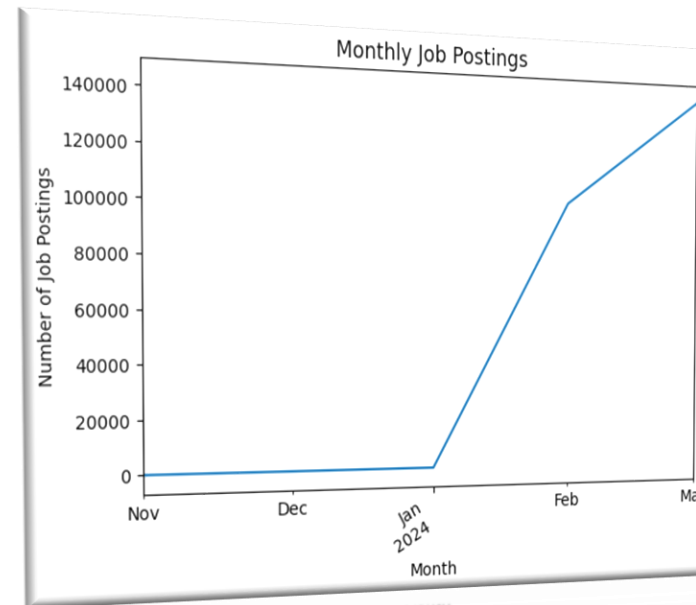
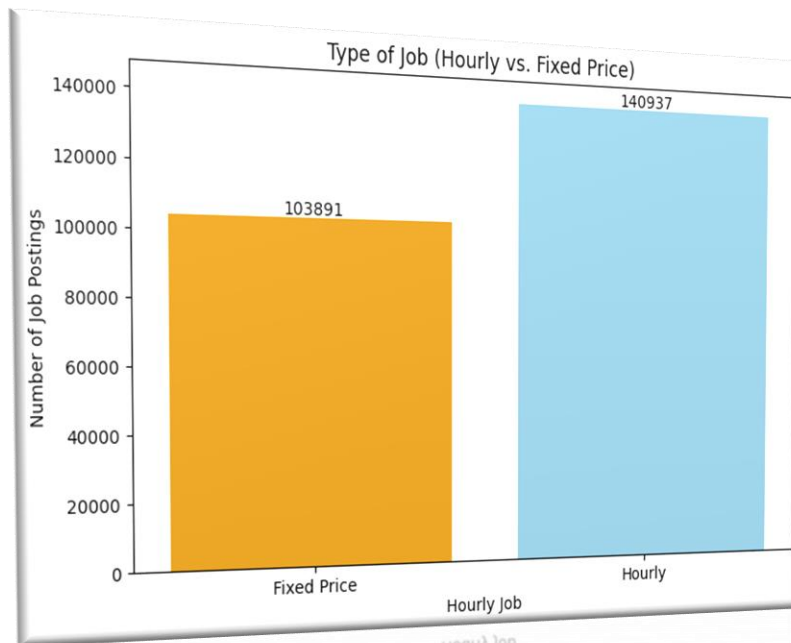
- Extracted keywords from **job titles** to find trending skills
- Found repeated terms mostly in **tech, writing & digital marketing**
- Currently rule-based; can improve with **NLP** or **embeddings**



Top Job-Posting Countries

- Grouped jobs by country to find where most roles come from
- Top 5 Countries:**
 - us USA – 99,834 jobs
 - gb UK – 19,149
 - in India – 15,835
 - au Australia – 12,618
 - ca Canada – 11,658
- us USA clearly dominates the platform

Task 1: Key Insights & Early Prototype



1

Hourly vs Fixed Price Jobs

- ☐ Compared job types based on **payment model**
- ☒ **Fixed-price jobs** were slightly more common than hourly
- ☐ Shows clients prefer **pay-per-project** structure

2

Budget Distribution

- ☐ Most job budgets are **low to moderate**
- ☐ A few **very high budgets** created skewed visuals
- ☐ Plan to apply **outlier caps** for better clarity next time

3

Monthly Posting Trend

- ☐ Grouped job posts by **month** to see job flow
- ☐ Created line charts to show **posting volume changes**
- ☐ Data is limited to 2 months, so trend scope is small

4

Basic Job Recommender (Prototype)

- ☐ Built a simple system to **match skills to jobs** by title
- ☐ Example: Input "Data Engineer" → returns all matching job titles
- ☐ Early version worked well; can be improved using **NLP techniques**



✓ Missing Value Cleanup

- ➡ Result: No missing values left in the dataset! 😊

Keyword Extraction from Titles

- Merged similar terms (e.g. "need" + "needed", "design" + "designer") for cleaner insights

- website, developer, expert, video, etc.

Word Cloud Snapshot

Offered a quick view of in-demand freelancing roles  

```

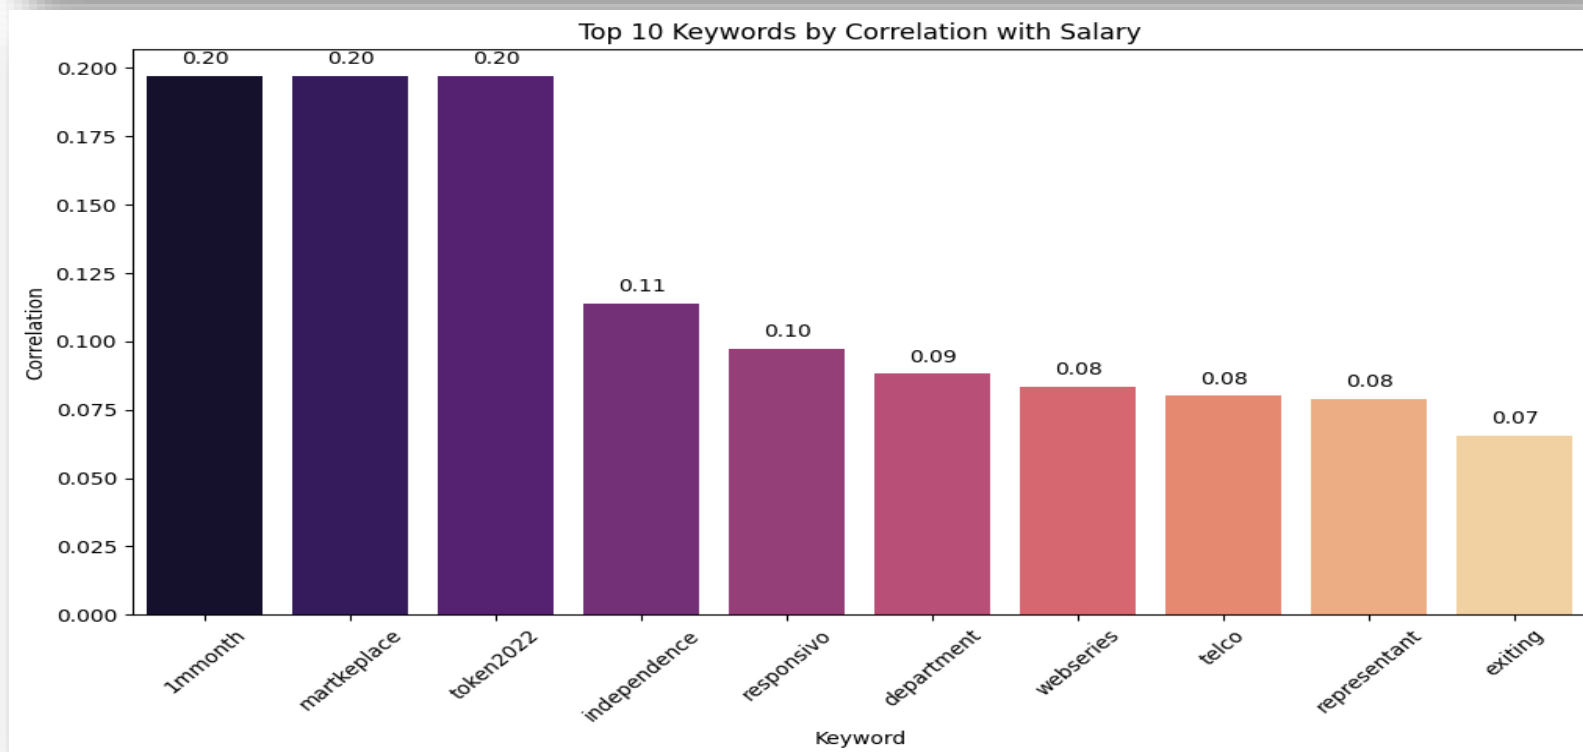

1 # Lemmatizing and Extracting Keywords from Job Titles
2 lemmatizer = WordNetLemmatizer() # Initialize the lemmatizer
3
4 # Function to lemmatize a single keyword
5 def lemmatize_keyword(keyword):
6     return lemmatizer.lemmatize(keyword)
7
8 # Function to extract and lemmatize keywords from a job title
9 def extract_keywords(title):
10     # Lowercase the title and remove punctuation
11     title = re.sub(r'[^\\w\\s]', '', title.lower())
12     # Tokenize the cleaned title into words
13     tokens = word_tokenize(title)
14     # Define English stop words to exclude common words
15     stop_words = set(stopwords.words('english'))
16     # Lemmatize tokens and filter out stop words
17     keywords = [lemmatize_keyword(token) for token in tokens if token not in stop_words]
18     return keywords
19
20 # Extract keywords from each job title and store them in a new column
21 df['keywords'] = df['title'].apply(extract_keywords)
22
23 # Flatten the list of keyword lists into a single list of keywords
24 keywords_flat = [keyword for sublist in df['keywords'] for keyword in sublist]
25
26 # Count how often each keyword appears
27 keyword_counts = Counter(keywords_flat)
28
29 # Get the top 10 most common keywords and their frequencies
30 top_10_keywords = keyword_counts.most_common(10)
31
32 # Convert the top keywords into a DataFrame for easier viewing
33 top_10_keywords_df = pd.DataFrame(top_10_keywords, columns=['Keyword', 'Frequency'])
34
35 # Display the top 10 keywords
36 top_10_keywords

```

[illegible]



Task 1 – Salary Insights from Keywords



01



Top Keywords by Avg. Salary

- Calculated average **budget per keyword**
- Found that **highest-paying keywords \neq most frequent**
 - Helped spot **niche roles** with **higher value**

02



Keyword-Salary Correlation

- Used **MultiLabelBinarizer** to create binary matrix for keywords
- Applied **Pearson correlation** between keyword presence & budget
- Found keywords with **strong positive salary correlation**
 - Indicates potential **premium skills** worth focusing on

03



Key Takeaways

- Keywords like **"developer"**, **"design"**, and **"expert"** are popular
- But **niche keywords** (less frequent) often pay more
 - Shows **low-supply, high-demand** opportunities in the freelance space



Task 2 – Identifying Emerging Job Categories

Classifying Job Titles into Categories Based on Keywords

```
1 def categorize_job(title):
2     # Loop through each category and its associated keywords
3     for category, keywords in categories.items():
4         # Check if any keyword is present in the job title (case-insensitive)
5         if any(keyword in title.lower() for keyword in keywords):
6             return category # Return the matching category
7     return 'Other' # If no keywords match, assign category as 'Other'
8
9 # Apply the categorize_job function to the 'title' column
10 df['category'] = df['title'].apply(categorize_job)
11
12 # Print the first 5 job titles with their assigned categories to check results
13 df[['title', 'category']].head(5)
```

	title	category
0	Experienced Media Buyer For Solar Pannel and R...	Other
1	Full Stack Developer	Software Development
2	SMMA Bubble App	Other
3	Talent Hunter Specialized in Marketing	Marketing
4	Data Engineer	Engineering

Calculate growth rate and identify emerging categories & Sort categories by growth rate to identify emerging categories

```
1 # Calculate month-over-month percentage change in job postings per category
2 growth_rate = category_trends.pct_change().fillna(0)
3
4 # Calculate the average growth rate for each category over the entire period
5 average_growth = growth_rate.mean()
6
7 # Sort categories by their average growth rate in descending order
8 emerging_categories = average_growth.sort_values(ascending=False)
9
10 # Print the sorted categories with their average growth rates
11 print("Categories sorted by average monthly growth rate:")
12 print(emerging_categories)
```

```
Categories sorted by average monthly growth rate:
category
Administration    inf
Construction       inf
Consulting         inf
Customer Service   inf
Data Science       inf
Design            inf
Education         inf
Engineering        inf
Finance           inf
Healthcare        inf
Hospitality       inf
Human Resources   inf
IT Support        inf
Legal            inf
Manufacturing     inf
Marketing         inf
Operations        inf
Product Management inf
Retail           inf
Project Management inf
Real Estate      inf
Research         inf
Software Development inf
Sales           inf
Writing and Editing inf
Other           261.991489
dtype: float64
```

01




Data Cleanup & Setup

- ☐ Converted published_date to datetime
- ☐ Created a new year_month column for monthly trend analysis 

02



Job Categorization

- ☐ Mapped job titles into categories like:
 - Data Science, Design, Consulting, Healthcare, etc.
- ☐ Used custom keyword matching
- ☐ Unmatched titles were labeled as 'Other' 

03




Monthly Job Posting Trends

- ☐ Used .groupby() + .unstack() to reshape data for time-series analysis
- ☐ Tracked job volume by category across months

04



Growth Rate Calculation

- ☐ Calculated month-over-month % change in job counts per category
- ☐ Averaged growth to identify fastest-growing categories 

Key Insights & Emerging Categories

STEP 01

★ Top Emerging Categories (by Avg. Growth)

- 1 Administration | 2 Construction | 3 Consulting → inf ⚠
 - 4 Customer Service | 5 Data Science → inf ⚠
- ⚠ inf = No jobs earlier, new now → indicates strong upward trend

STEP 02

Visualization

- ☐ Line chart created to track top 5 growing categories
- ☐ Showed fields like Data Science and Consulting rising steadily 🔍

STEP 03

Caveats

- ☐ Some spikes due to new categories appearing for the first time
- ☐ Not all indicate consistent demand yet - but useful early indicators 🔍

STEP 04

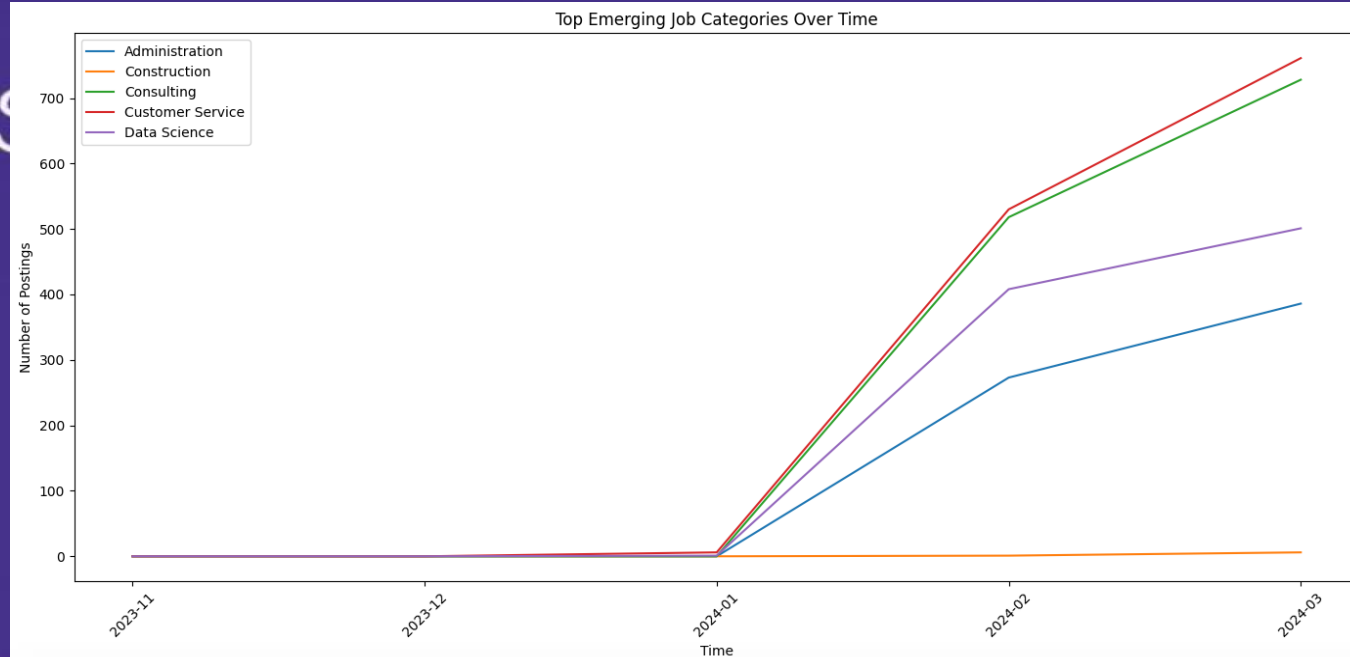
Final Step

- ☐ Saved the cleaned and categorized dataset for future use
 - 📁 /job_data.csv

STEP 05

✅ Summary

- ☐ Discovered emerging job categories by tracking monthly trends
- ☐ Fields like Data Science, Customer Service, Consulting show high momentum
- ☐ Useful for platforms and recruiters to focus on upcoming demand areas 🧠



Tracking Top Emerging Job Categories Over Time

```
1 import matplotlib.pyplot as plt
2
3 # Set figure size for better visibility
4 plt.figure(figsize=(14, 7))
5
6 # Loop through the top 5 emerging categories based on average growth rate
7 for category in emerging_categories.index[:5]:
8     # Plot the number of postings over time for each category
9     plt.plot(category_trends.index.astype(str), category_trends[category], label=category)
10
11 # Add title and axis labels
12 plt.title('Top Emerging Job Categories Over Time')
13 plt.xlabel('Time')
14 plt.ylabel('Number of Postings')
15
16 # Add legend to identify categories
17 plt.legend()
18
19 # Rotate x-axis labels for better readability
20 plt.xticks(rotation=45)
21
22 # Adjust layout to prevent clipping of labels
23 plt.tight_layout()
24
25 # Display the plot
26 plt.show()
27
28 # Print the top 10 emerging categories with their average growth rates
29 print("Emerging Job Categories with Average Growth Rates:")
30 print(emerging_categories.head(10))
```

Saving Job data to CSV file

```
[ ] 1 # folder to save file
2   folder_path = '/content/drive/MyDrive/Job Market Analysis & Recommendation System'
3
4   # csv file name
5   file_name = 'job_data.csv'
6
7   # Combine folder path and file name to create full file path
8   csv_filepath = f'{folder_path}/{file_name}'
9
10  # Save DataFrame to the specified folder with the given file name
11  df.to_csv(csv_filepath, index=False)
12
13  print(f"Data saved as {csv_filepath} with Filename: {file_name}")
```

Data saved as /content/drive/MyDrive/Job Market Analysis & Recommendation System/job_data.csv with Filename: job_data.csv

Emerging Job Categories with Average Growth Rates:

category	
Administration	inf
Construction	inf
Consulting	inf
Customer Service	inf
Data Science	inf
Design	inf
Education	inf
Engineering	inf
Finance	inf
Healthcare	inf
dtype:	float64

Task 3: Predicting High-Demand Job Roles | Forecasting & Insights

1



Data Preparation

- ❑ Extracted key fields: **title**, **category**, **month**
- ❑ Created a **target variable** for monthly job demand
- ❑ Applied **scaling, encoding & dimensionality reduction** for model training

2



Modeling

- ❑ Trained a **Gradient Boosting Regressor**
- ❑ Used **MAE, MSE, and R²** for evaluation
 - ⚠️ Model performance was low due to **data sparsity & noise**

3



Trend Analysis

- ❑ Analyzed job posting changes month-by-month
- ❑ Identified **Top 10 Emerging Job Titles** based on growth:
 - ✅ Social Media Manager | ✅ Video Editor | ✅ Logo Design | ✅ Graphic Designer
 - ✅ Virtual Assistant | ✅ Content Writer | ✅ YouTube Manager | ✅ Personal Assistant
 - ✅ Web Designer | ✅ Bookkeeper

4



Time Series Forecasting

- ❑ Applied **exponential smoothing** to predict future trends
- ❑ Generated 12-month forecasts for each category
 - 📌 Example: **Data Science** shows steady upward trend
- ❑ Plotted forecasts to visualize **emerging or seasonal demand**

5



Key Insights

- 👤💻 Creative + remote-friendly roles like
 - **Social Media Manager** and **Virtual Assistant** are rising fast
- 📄 Steady demand remains for
 - **Web Designers, Content Writers, and Bookkeepers**
- 📊 Forecasts give job seekers and platforms
 - A clear view of **what's trending and where to focus** 🧠

Define Features, Target, and Data Preprocessing Pipeline

```
1 from sklearn.preprocessing import StandardScaler, OneHotEncoder
2 from sklearn.compose import ColumnTransformer
3
4 # Create 'demand' target by counting job titles per 'year_month'
5 job_data['demand'] = job_data.groupby('year_month')['title'].transform('count')
6
7 # Define features by dropping columns not used as inputs
8 features = job_data.drop(columns=['demand', 'published_date', 'year_month', 'link'])
9
10 # Define target variable
11 target = job_data['demand']
12
13 # Specify numerical and categorical feature columns
14 numerical_features = ['hourly_low', 'hourly_high', 'budget']
15 categorical_features = ['title', 'country', 'keywords', 'category', 'is_hourly']
16
17 # Define transformers for numerical and categorical features
18 numerical_transformer = StandardScaler()
19 categorical_transformer = OneHotEncoder(handle_unknown='ignore', sparse_output=True)
20
21 # Create ColumnTransformer to apply transformations appropriately
22 preprocessor = ColumnTransformer(
23     transformers=[
24         ('num', numerical_transformer, numerical_features),
25         ('cat', categorical_transformer, categorical_features)
26     ],
27     sparse_threshold=0.3 # output sparse matrix if >30% zeros
28 )
29
30 print("Preprocessing pipeline created successfully.")
```

Preprocessing pipeline created successfully.

Gradient Boosting Regressor


```
1 from sklearn.ensemble import GradientBoostingRegressor
2 from sklearn.metrics import mean_absolute_error
3 import matplotlib.pyplot as plt
4
5 # Initialize the Gradient Boosting Regressor
6 gb_model = GradientBoostingRegressor(n_estimators=100, random_state=42)
7
8 # Train the model on the training data
9 gb_model.fit(X_train, y_train)
10
11 # Predict on the test set
12 y_pred_gb = gb_model.predict(X_test)
13
14 # Evaluate the model using Mean Absolute Error
15 mae_gb = mean_absolute_error(y_test, y_pred_gb)
16 print(f'Mean Absolute Error with Gradient Boosting: {mae_gb}')
```

Mean Absolute Error with Gradient Boosting: 19814.832581437848

title		
Social Media Manager		61.25
Video Editor		50.50
Logo Design		50.25
Graphic Designer		43.00
Virtual Assistant		41.50
Logo design		25.00
Logo Designer		22.00
Lead Generation		20.75
Full Stack Developer		20.00
Appointment Setter		18.00
dtype: float64		

Prepare Time Series Data for Job Postings by Category

```
1 import pandas as pd
2
3 # Convert the 'year_month' column to datetime
4 job_data['year_month'] = pd.to_datetime(job_data['year_month'])
5
6 # Group by date and category, and count the number of postings
7 time_series_data = job_data.groupby(['year_month', 'category']).size().unstack().fillna(0)
8
9 # Display the prepared data
10 print(time_series_data.head())
```





Task 4 – Exploring Hourly Rate Differences Across Countries

01 Objective

To explore how **freelance hourly rates vary** across countries & identify **high-paying regions** 💰🌍

02 What I Did

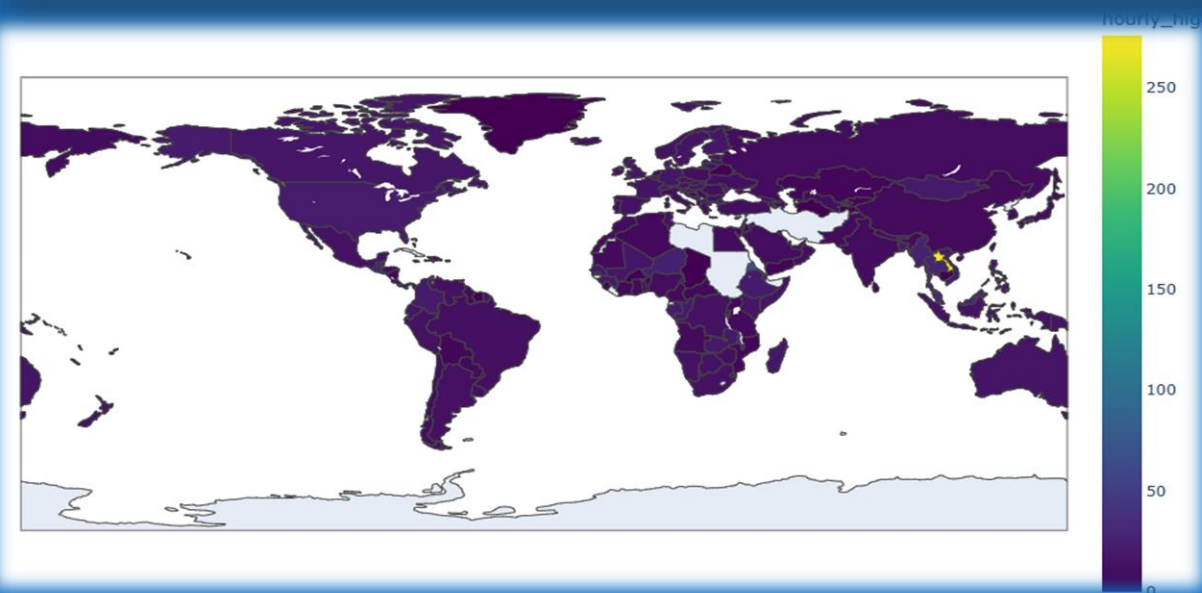
- ❑ Grouped jobs by **country** and calculated average of hourly_high
- ❑ Built a **choropleth map** using Plotly Express to visualize global pay trends
- ❑ Used **Viridis colour scale** for contrast and customized layout for clarity 🌈
- ❑ Created a description field combining **job title + country** for future keyword-based tasks

03 Key Findings

- 🌟 Countries like **Vietnam** and **Philippines** showed **surprisingly high average rates**
- 🌍 Lower rates were common in **India, Africa, and South America**
- 👤 Helped highlight **regional pay disparities** and where to target high-paying remote work
- 📍 Interactive map made it easy to explore hourly rates country-wise by hovering

Average Hourly Rates by Country

```
1 import plotly.express as px
2
3 # Calculate average hourly rates by country
4 avg_hourly_rates = job_data.groupby('country')['hourly_high'].mean().reset_index()
5
6 # Plot choropleth map
7 fig = px.choropleth(
8     avg_hourly_rates,
9     locations='country',
10    locationmode='country names',
11    color='hourly_high',
12    hover_name='country',
13    title='Average Hourly Rates by Country (Values can be seen on hover)',
14    color_continuous_scale='Viridis'
15 )
16
17 # Increase figure size to better fit screen width
18 fig.update_layout(
19     width=1000,
20     height=700
21 )
22
23 fig.show()
```



```
[ ] 1 # Show full column content
2     pd.set_option('display.max_colwidth', None)
3
4 # Create a new 'description' column by concatenating 'title' and 'country' with a space in between
5 job_data['description'] = job_data['title'] + ' ' + job_data['country']
6
7 # Print the first 5 rows to verify the new column
8 job_data[['title', 'country', 'description']].head()
```

	title	country	description
0	Experienced Media Buyer For Solar Pannel and Roofing installation companies.	Other	Experienced Media Buyer For Solar Pannel and Roofing installation companies. Other
1	Full Stack Developer	United States	Full Stack Developer United States
2	SMMA Bubble App	United States	SMMA Bubble App United States
3	Talent Hunter Specialized in Marketing	United States	Talent Hunter Specialized in Marketing United States
4	Data Engineer	India	Data Engineer India



Task 5 – Job Similarity Recommendation Using TF-IDF & Cosine Similarity

Extracting Features from Job Descriptions Using TF-IDF Vectorization

```
1 from sklearn.feature_extraction.text import TfidfVectorizer
2
3 # Initialize TF-IDF Vectorizer with English stop words removed
4 vectorizer = TfidfVectorizer(stop_words='english')
5
6 # Fit the vectorizer on the 'description' column and transform the text data into TF-IDF features
7 X = vectorizer.fit_transform(job_data['description'])
8
9 # X is a sparse matrix of shape (number_of_samples, number_of_features)
10 print(f"TF-IDF matrix shape: {X.shape}")
11
12 # Optionally, get the feature names (words)
13 feature_names = vectorizer.get_feature_names_out()
14 print(f"Number of unique features: {len(feature_names)}")
```

TF-IDF matrix shape: (244827, 38928)
Number of unique features: 38928

Building a Nearest Neighbors Model with Cosine Similarity for Job Descriptions

```
[ ] 1 from sklearn.neighbors import NearestNeighbors
2
3 # Initialize Nearest Neighbors model with cosine distance and brute-force algorithm
4 model = NearestNeighbors(metric='cosine', algorithm='brute')
5
6 # Fit the model on the TF-IDF feature matrix
7 model.fit(X)
```

NearestNeighbors
NearestNeighbors(algorithm='brute', metric='cosine')

```
1 # Find the 5 nearest neighbors for the first job description
2 distances, indices = model.kneighbors(X[0], n_neighbors=5)
3
4 # Print formatted output
5 print("Top similar job descriptions to the first job:\n")
6
7 for rank, (idx, dist) in enumerate(zip(indices[0], distances[0]), start=1):
8     similarity = 1 - dist # cosine similarity = 1 - cosine distance
9     description = job_data.iloc[idx]['description']
10    print(f"{rank}. Index: {idx}, Similarity: {similarity:.3f}")
11    print(f"    Description: {description}\n")
```

Top similar job descriptions to the first job:

1. Index: 0, Similarity: 1.000
Description: Experienced Media Buyer For Solar Pannel and Roofing installation companies. Other
2. Index: 221025, Similarity: 0.545
Description: Experienced Facebook and YouTube Ads Media Buyer for Roofing companies Morocco
3. Index: 190320, Similarity: 0.523
Description: Database of Solar Installation companies in US United Kingdom
4. Index: 160365, Similarity: 0.514
Description: Media Buyer Needed to generate leads for Solar pannel company (Google + Meta Ads) Belgium
5. Index: 114504, Similarity: 0.511
Description: Expert Media Buyer For Roofing United States




Text Vectorization Using TF-IDF

- Cleaned job descriptions and applied **TF-IDF** to highlight important terms
- Final matrix: **244,827 jobs × 38,928 features** — showcasing rich textual variety
- Removed common filler words (e.g., "the", "and") to focus on meaningful terms



Similarity Matching with Nearest Neighbors

- Used **cosine similarity** with Nearest Neighbors to compare job descriptions
- Retrieved **top 5 most similar jobs** for any given listing
- Lower cosine distance = higher similarity**
- Added a custom description field (**title + country**) for improved context



Model Saving for Reuse

- Saved both **TF-IDF vectorizer** and **Nearest Neighbors model** using pickle
- Enabled **reusability** without retraining — ideal for deployment or integration



Performance & Output Insights

- Recommendations were **accurate and relevant** — cosine scores around **0.51–0.54**
- Similar jobs were found even across **different countries** with contextual overlap
- Model is **fast, scalable, and easy to interpret**, even with **200K+ jobs**
- Extracted full **country list** to support **region-specific recommendations**



Task 6 – Tracking Monthly Changes in Job Market Demand

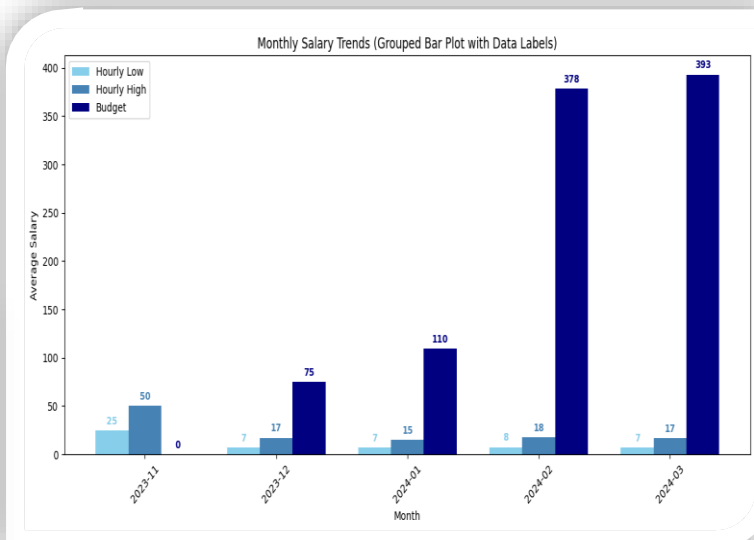
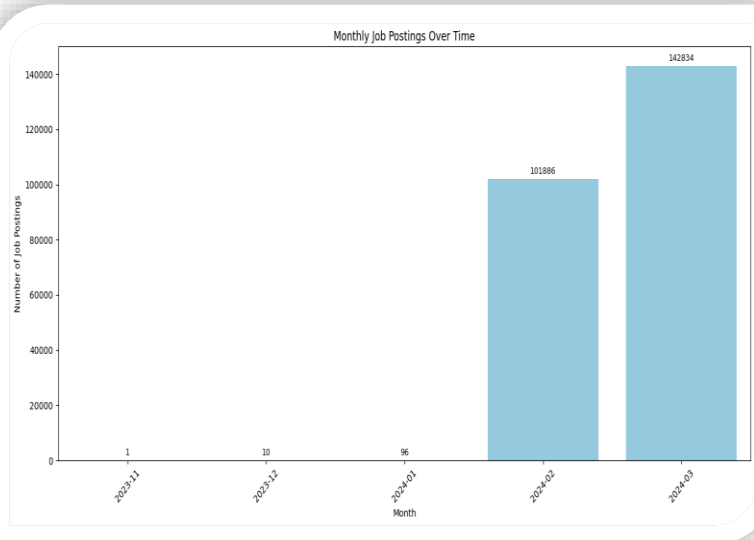
Monthly Trends in Job Postings, Hourly Rates, and Budgets

```
1 # Group by 'year_month' to get the count of job postings per month
2 monthly_job_postings = job_data1.groupby('year_month').size().reset_index(name='job_postings')
3
4 # Calculate average hourly_low, hourly_high, and budget per month
5 monthly_salary_trends = job_data1.groupby('year_month').agg({
6     'hourly_low': 'mean',
7     'hourly_high': 'mean',
8     'budget': 'mean'
9 }).reset_index()
10
11 # Merge the counts and average salary/budget trends into one DataFrame
12 monthly_trends = pd.merge(monthly_job_postings, monthly_salary_trends, on='year_month')
13
14 # Display the first few rows of the combined monthly trends
15 print(monthly_trends.head())
```

	year_month	job_postings	hourly_low	hourly_high	budget
0	2023-11	1	25.000000	50.000000	0.000000
1	2023-12	10	6.900000	17.200000	75.000000
2	2024-01	96	6.739583	15.083333	109.656250
3	2024-02	101886	7.523045	17.520601	378.296832
4	2024-03	142834	7.062233	16.514786	393.039976

Tracking Monthly Job Postings Over Time

```
1 import matplotlib.pyplot as plt
2 import seaborn as sns
3
4 plt.figure(figsize=(14, 7))
5 barplot = sns.barplot(data=monthly_trends, x='year_month', y='job_postings', color='skyblue')
6
7 # Add data labels on top of each bar
8 for p in barplot.patches:
9     height = p.get_height()
10    barplot.annotate(f'{int(height)}',
11                    (p.get_x() + p.get_width() / 2, height),
12                    ha='center', va='bottom',
13                    fontsize=9, color='black', xytext=(0, 3),
14                    textcoords='offset points')
15
16 plt.title('Monthly Job Postings Over Time')
17 plt.xlabel('Month')
18 plt.ylabel('Number of Job Postings')
19 plt.xticks(rotation=45)
20 plt.tight_layout()
21 plt.show()
```



01

Most In-Demand Job Roles

- Frequently posted roles include:
 - Data Scientist | Data Analyst | ML Engineer | Business Analyst
- Highlights a strong industry focus on **data-driven roles**

02

Overall Growth in Job Postings

- Observed a **steady rise in job volume** month-over-month
- Suggests rising demand for skilled professionals
- Driven by digital adoption, AI/ML integration & data-first strategies

03

High Demand for Technical Skills

- Most roles require:
 - Python / R | Machine Learning | SQL | Power BI / Tableau
- Shows strong demand for **analytical & programming expertise**.

04

Seasonal Fluctuations Noted

- Despite the growth, **up-and-down patterns** appear over months
- Likely caused by:
 - Quarterly hiring cycles | Budget approvals | Project kick-offs

05

Guidance for Job Seekers To Focus On

- In-demand roles & tools**
- Staying updated on job market trends
- Practicing real-world tech skills



Task 7 – Remote Work Trends Analysis Using Rolling Average & Forecasting

Monthly Analysis of Remote Job Postings Based on Keywords

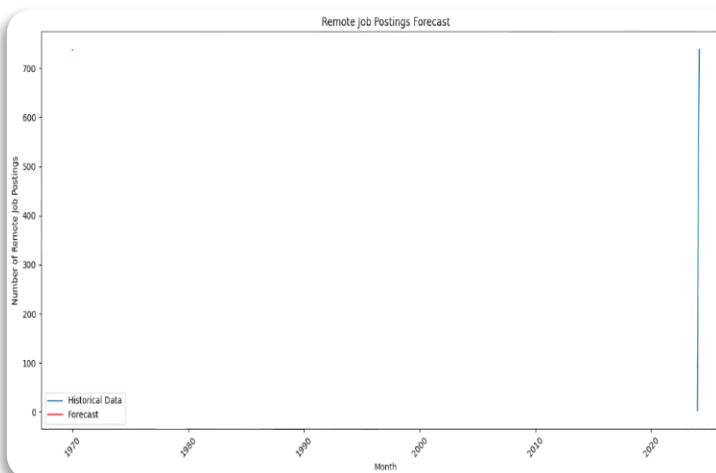
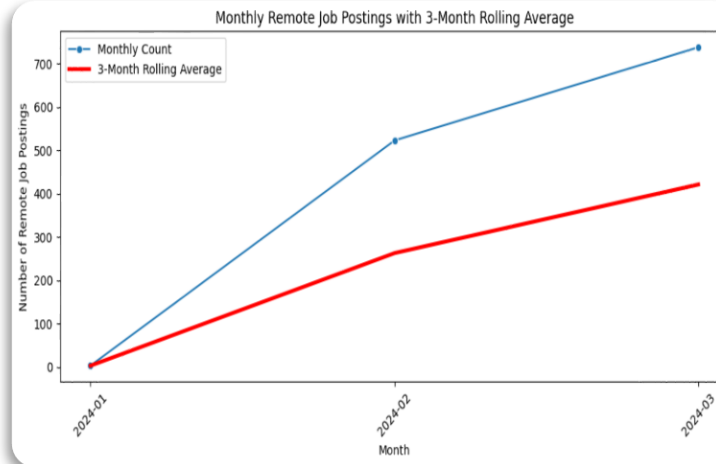
```
[ ] 1 # Fill missing values in 'keywords' column with empty strings
2 job_data1['keywords'] = job_data1['keywords'].fillna('')
3
4 # Filter job postings that mention 'remote' in the keywords (case-insensitive)
5 remote_jobs = job_data1[job_data1['keywords'].str.contains('remote', case=False, na=False)]
6
7 # Group by 'year_month' and count the number of remote job postings per month
8 monthly_remote_jobs = remote_jobs.groupby('year_month').size().reset_index(name='remote_job_postings')
9
10 # Convert 'year_month' to string type for easier plotting
11 monthly_remote_jobs['year_month'] = monthly_remote_jobs['year_month'].astype(str)
12
13 # Print the first few rows of the result
14 print(monthly_remote_jobs.head())
```

```
year_month  remote_job_postings
0    2024-01                    3
1    2024-02                   523
2    2024-03                   738
```

Forecasting Remote Job Posting Trends Using Exponential Smoothing

```
1 import pandas as pd
2 from statsmodels.tsa.holtwinters import ExponentialSmoothing
3
4 # Convert 'year_month' to datetime timestamp for modeling
5 monthly_remote_jobs['year_month'] = pd.PeriodIndex(monthly_remote_jobs['year_month'], freq='M').to_timestamp()
6
7 # Calculate initial seasonal values using a 12-month rolling mean
8 seasonal_initial = monthly_remote_jobs['remote_job_postings'].rolling(window=12, min_periods=1).mean()
9
10 # Set initial level and seasonal components for the model
11 initial_level = monthly_remote_jobs['remote_job_postings'].iloc[0]
12 initial_seasonal = seasonal_initial.iloc[0]
13
14 # Define and fit the Holt-Winters Exponential Smoothing model with additive seasonality
15 model = ExponentialSmoothing(
16     monthly_remote_jobs['remote_job_postings'],
17     seasonal='add',
18     seasonal_periods=12,
19     initialization_method='known',
20     initial_level=initial_level,
21     initial_seasonal=initial_seasonal
22 )
23 fit = model.fit()
24
25 # Forecast the next 12 months
26 forecast = fit.forecast(12)
27
28 # Print the forecasted values
29 print(forecast)
```

```
3    736.912000
4    736.912000
5    736.912000
6    736.912000
7    736.912000
8    736.912000
9    736.912000
10   736.912000
11   736.912000
12   736.911700
13   736.963999
14   736.912000
dtype: float64
```



Visualization of Monthly Trends in Remote Job Postings with 3-Month Rolling Average

```
1 import matplotlib.pyplot as plt
2 import seaborn as sns
3
4 # Calculate 3-month rolling average to smooth data
5 monthly_remote_jobs['rolling_avg'] = monthly_remote_jobs['remote_job_postings'].rolling(window=3, min_periods=1).mean()
6
7 plt.figure(figsize=(10, 5)) # Set plot size
8
9 # Plot original monthly remote job postings
10 sns.lineplot(data=monthly_remote_jobs, x='year_month', y='remote_job_postings', label='Monthly Count', marker='o')
11
12 # Plot rolling average with emphasis
13 sns.lineplot(data=monthly_remote_jobs, x='year_month', y='rolling_avg', label='3-Month Rolling Average', color='red', linewidth=3)
14
15 plt.title('Monthly Remote Job Postings with 3-Month Rolling Average')
16 plt.xlabel('Month')
17 plt.ylabel('Number of Remote Job Postings')
18 plt.xticks(rotation=45) # Rotate x-axis labels for readability
19 plt.legend()
20 plt.tight_layout()
21 plt.savefig('remote_job_postings_rolling_avg.png')
22 plt.show()
```

01

Surge in Remote Job Postings

- Remote jobs increased from 3 in Jan to 738 in Mar 2024
- Proves remote work is quickly rising in popularity

02

Rolling Average Reveals Consistent Growth

- Applied a 3-month rolling average to smooth out spikes
- Clearly shows a steady upward momentum in demand

03

Forecast Indicates Stable Future Demand

- Used Exponential Smoothing to predict next 12 months
- Remote job postings likely to stay above 736/month

04

Shift in Job Market Behaviour

- Compared to early months, demand shows rapid change
- Remote flexibility is now a mainstream hiring norm



Task 8 – Analyzing Overall Job Market Trends & Forecasts

Analyze Monthly Job Postings by Year and Month

```
1 import pandas as pd
2 import matplotlib.pyplot as plt
3 import seaborn as sns
4 from statsmodels.tsa.arima.model import ARIMA

1 # Group the job data by 'year_month' to count the number of job postings per month
2 monthly_job_postings = job_data1.groupby('year_month').size().reset_index(name='job_postings')
3
4 # Print the first few rows to verify the grouping and counts
5 print("Grouped monthly job postings (first 5 rows):")
6 print(monthly_job_postings.head())
7
8 # Convert 'year_month' to string type for plotting or display purposes
9 monthly_job_postings['year_month'] = monthly_job_postings['year_month'].astype(str)
10
11 # Print data types to confirm conversion
12 print("\nData types after conversion:")
13 print(monthly_job_postings.dtypes)
```


```
Grouped monthly job postings (first 5 rows):
  year_month  job_postings
0  2023-11             1
1  2023-12            10
2  2024-01            96
3  2024-02          101886
4  2024-03          142834

Data types after conversion:
year_month      object
job_postings    int64
dtype: object
```

Forecast Future Job Postings Using ARIMA Model

```
1 from statsmodels.tsa.arima.model import ARIMA
2 import pandas as pd
3
4 # Sort data by date
5 monthly_job_postings = monthly_job_postings.sort_values('year_month')
6
7 # Fit ARIMA model
8 model = ARIMA(monthly_job_postings['job_postings'], order=(1, 1, 1))
9 fit_model = model.fit()
10
11 # Forecast next 12 months
12 forecast_steps = 12
13 forecast = fit_model.forecast(steps=forecast_steps)
14
15 # Prepare index for forecasted months
16 last_date = pd.to_datetime(monthly_job_postings['year_month'].iloc[-1])
17 forecast_index = pd.date_range(start=last_date + pd.offsets.MonthBegin(),
18                               periods=forecast_steps, freq='MS').strftime('%Y-%m')
19
20 # Combine forecasted values with their dates
21 forecast_df = pd.DataFrame({'year_month': forecast_index, 'forecasted_job_postings': forecast.values})
22
23 print("Forecasted job postings for next 12 months:\n")
24 print(forecast_df)
```

Forecasted job postings for next 12 months:		
	year_month	forecasted_job_postings
0	2024-04	159319.147364
1	2024-05	165955.874347
2	2024-06	168627.742791
3	2024-07	169703.405690
4	2024-08	170136.454957
5	2024-09	170310.795513
6	2024-10	170380.982976
7	2024-11	170409.239625
8	2024-12	170420.615421
9	2025-01	170425.195183
10	2025-02	170427.038941
11	2025-03	170427.781217





Huge Spike in Q1 2024

-  Job postings surged in **Feb (101K)** and **Mar (143K)**
-  Indicates **high hiring demand** after a quiet quarter



Sharp Recovery from Zero Activity

-  Nov 2023 to Jan 2024 had **near-zero postings**
-  Feb 2024 brought a strong **market rebound** — likely seasonal






Future Forecast – Growth Then Stability

-  ARIMA model shows growth until **July 2024 (~170K)**
-  After that, **job volume plateaus**, showing **market stabilization**




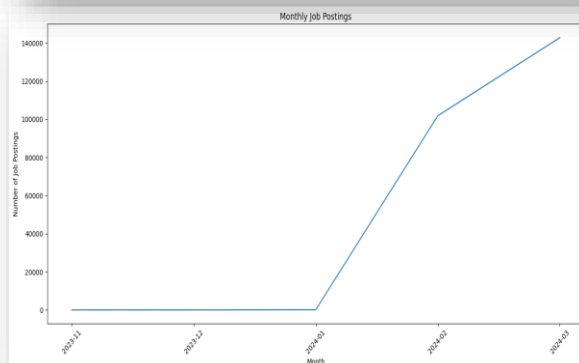
Forecast Shows High Confidence

-  Forecast trend is **smooth, not volatile**
-  Suggests market is entering a **mature, stable phase**



Seasonality Influences Hiring Patterns

-  Spike aligns with **post-holiday budgets & hiring cycles**
-  Future trend looks like **steady, strategic hiring**



AI-Powered Job Role Recommender

Get top job suggestions based on your description, with filters.

Enter Job Description

Example: Facebook and YouTube Ads Media Buyer for Roofing companies

Describe the job role you're looking for...

Experience Level(s)

Fresher ×

Experienced ×

Country/Countries

Austria ×

Bermuda ×

Job Type(s)


Remote ×

Full-Time ×

Recommend Jobs

Reset All

Top 5 Available Countries



Building the Streamlit-Based Job Recommendation System

End-to-End App Functionality

- Built an **AI-powered job recommender app** using **Streamlit**
- Accepts a **user-entered job description** and returns **top matching jobs**
- Allows filters for:
 - **Experience level** (Fresher / Experienced)
 - **Country**
 - **Job type** (Remote, On-site, Freelance, etc.)

Recommendation Logic

- Vectorized input using a **pre-trained TF-IDF model**
- Compared against 240K+ job descriptions using **cosine similarity**
- Returned **top 10 most similar roles** ranked by match percentage
- Used **highlighted keywords** to show why a job was recommended
- Used a **Nearest Neighbors model**, loaded via pickle for efficiency

Data Handling & Preprocessing




- Cleaned the dataset by filling missing values and creating extra fields
- Merged **job title + country** into a single description field for better context
- Handled missing fields like **experience**, **job_type** and **keywords** smoothly
- App shows an **error and stops** if the model or data isn't available

Output Features & User Experience of the Streamlit App

 **Live Streamlit App:** [App Link](#)

 **All About My App:** [Video Link](#)

Output Visualization

- ☐ Displays job matches with:
 - ☐  **Highlighted titles**
 - ☐  **Location, Experience, Type, and Posting Date**
 - ☐  **Similarity score in %**
- ☐ Results are styled for **readability** and **scrollable exploration**



01

Top Country Insights

- ☐  Extracted and visualized **top 5 countries** with most matching jobs
- ☐  Used **Matplotlib bar chart** for a clean and clear overview





02

PDF Report Generation

- ☐ Added a feature to **download results as a PDF**:
- ☐  Job title, location, experience, job type
- ☐  Keywords and similarity %
- ☐ Styled with **FPDF** for clear formatting and readable output

03

Final Highlights

- ☐ ☒  Real-time filtering, smart recommendations, and export-ready results
- ☐  Fast, interpretable, and scalable solution for large job datasets
- ☐  All resources (model, vectorizer, dataset) loaded efficiently using caching
- ☐  Offers a **complete job search experience** within a single app

04

Enter Job Description

Example: Facebook and YouTube Ads Media Buyer for Roofing companies

Facebook and YouTube Ads Media Buyer for Roofing companies

Experience Level(s)

Fresher ×

Experienced ×

Country/Countries

Austria ×

Bermuda ×

Job Type(s)

Remote ×

Full-Time ×

Recommend Jobs

Reset All

No jobs found in selected country/countries, but available in: United States, Morocco, Other, India

Top Matching Job Roles:

Experienced Facebook and YouTube Ads Media Buyer for Roofing companies

Location: Morocco Date: 2024-03-18 Experience: Not Specified Type: Unknown
Keywords: ['experienced', 'facebook', 'youtube', 'ad', 'medium', 'buyer', 'roofing', 'company']
Similarity: 88.3%

Media Buyer for YouTube and Facebook Ads

Location: Other Date: 2024-02-19 Experience: Not Specified Type: Unknown
Keywords: ['medium', 'buyer', 'youtube', 'facebook', 'ad']
Similarity: 71.73%

Expert Media Buyer For Roofing

Location: United States Date: 2024-03-03 Experience: Not Specified Type: Unknown
Keywords: ['expert', 'medium', 'buyer', 'roofing']
Similarity: 68.64%

Media buyer (Facebook Ads)

Location: United States Date: 2024-02-20 Experience: Not Specified Type: Unknown
Keywords: ['medium', 'buyer', 'facebook', 'ad']
Similarity: 64.5%

Media Buyer for Facebook Ads

Location: United States Date: 2024-02-22 Experience: Not Specified Type: Unknown
Keywords: ['medium', 'buyer', 'facebook', 'ad']
Similarity: 64.5%

Facebook Ads Media Buyer

Location: United States Date: 2024-03-02 Experience: Not Specified Type: Unknown
Keywords: ['facebook', 'ad', 'medium', 'buyer']
Similarity: 64.5%



○ ○

Thank You!

Hope You find it helpful

